

Importing Intel Library

- We are running this intel patch in the beginning so all the sklearn libraries imported will have intel extension for faster processing

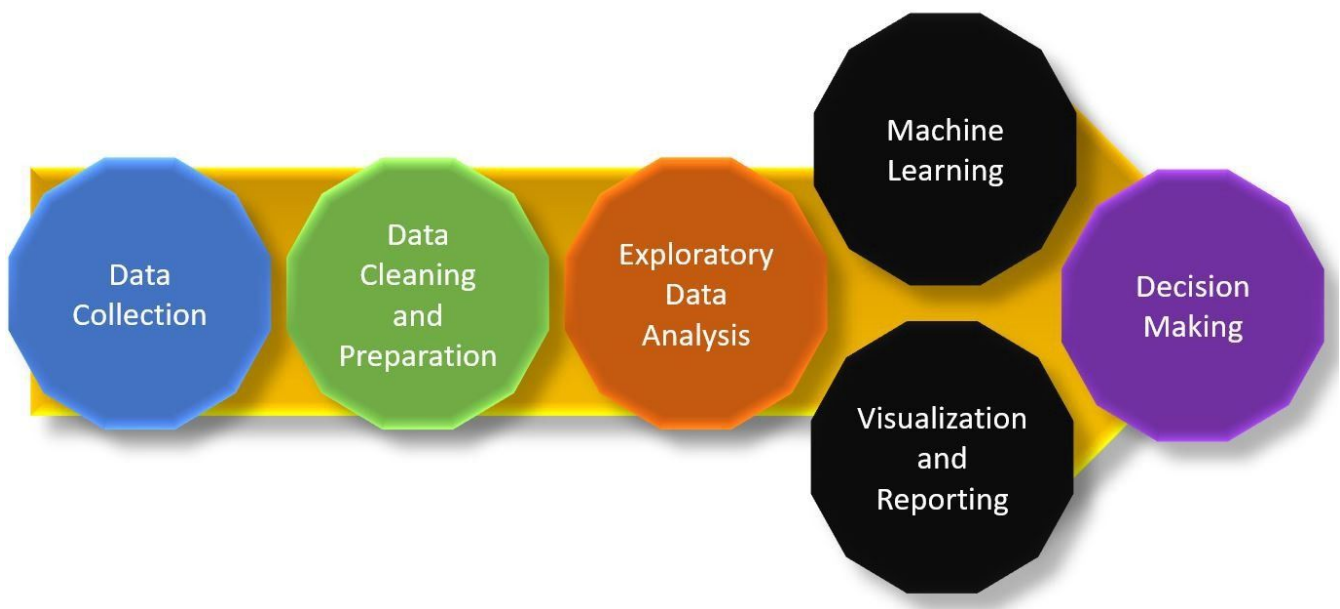


```
In [4]: 1 from sklearnex import patch_sklearn
        2 patch_sklearn()
```


Intel(R) Extension for Scikit-learn* enabled (<https://github.com/intel/scikit-learn-in-telex>)

Importing All the Libraries

```
In [5]: 1 import pandas as pd
        2 import numpy as np
        3 from numpy import mean
        4 from numpy import std
        5 import matplotlib.pyplot as plt
        6 import seaborn as sns
        7 from sklearn.metrics import f1_score
        8 from sklearn.linear_model import LogisticRegression
        9 from sklearn.model_selection import train_test_split
       10 from sklearn.metrics import mean_squared_error, r2_score
       11 import xgboost
       12 from sklearn import ensemble
       13 from sklearn.linear_model import LogisticRegression
       14 from sklearn.preprocessing import StandardScaler
       15 from xgboost import XGBClassifier
       16 from lightgbm import LGBMClassifier
       17 from catboost import CatBoostClassifier
       18 from sklearn.ensemble import GradientBoostingClassifier
       19 from sklearn.model_selection import cross_val_score
       20 from sklearn.model_selection import RepeatedStratifiedKFold
```



Loading the Dataset



```
In [6]: 1 data = pd.read_csv('dataset.csv', index_col = 0)
```

In [7]: 1 data.head(5)

Out[7]:

	pH	Iron	Nitrate	Chloride	Lead	Zinc	Color	Turbidity	Fluoride	Copper
Index										
0	8.332988	0.000083	8.605777	122.799772	3.713298e-52	3.434827	Colorless	0.022683	0.607283	0.14459
1	6.917863	0.000081	3.734167	227.029851	7.849262e-94	1.245317	Faint Yellow	0.019007	0.622874	0.43783
2	5.443762	0.020106	3.816994	230.995630	5.286616e-76	0.528280	Light Yellow	0.319956	0.423423	0.43158
3	7.955339	0.143988	8.224944	178.129940	3.997118e-176	4.027879	Near Colorless	0.166319	0.208454	0.23945
4	8.091909	0.002167	9.925788	186.540872	4.171069e-132	3.807511	Light Yellow	0.004867	0.222912	0.61657

5 rows × 23 columns

Data Cleansing

In [8]: 1 data.describe()

Out[8]:

	pH	Iron	Nitrate	Chloride	Lead	Zinc	Turbidity
count	5.840788e+06	5.917089e+06	5.851117e+06	5.781311e+06	5.929933e+06	5.800716e+06	5.907027e+06
mean	7.445373e+00	1.279027e-01	6.169970e+00	1.842970e+02	1.498336e-03	1.550255e+00	5.215093e-01
std	8.881665e-01	4.799915e-01	3.256667e+00	6.842828e+01	3.250641e-02	1.546368e+00	9.258807e-01
min	1.057113e+00	2.047587e-53	2.861727e-01	2.363919e+01	0.000000e+00	1.482707e-08	1.029712e-16
25%	6.894328e+00	9.992949e-06	3.973078e+00	1.381341e+02	1.500283e-122	4.148202e-01	3.872368e-02
50%	7.449564e+00	2.249640e-03	5.604051e+00	1.760178e+02	2.213625e-62	1.081818e+00	2.097680e-01
75%	8.014424e+00	5.455290e-02	7.672402e+00	2.179811e+02	3.592165e-27	2.230841e+00	6.249132e-01
max	1.291072e+01	1.935315e+01	9.639078e+01	1.507310e+03	5.844281e+00	2.836867e+01	2.371527e+01

```
In [9]: 1 data.isna().sum()
```

```
Out[9]: pH          116054
        Iron         39753
        Nitrate      105725
        Chloride     175531
        Lead         26909
        Zinc         156126
        Color        5739
        Turbidity    49815
        Fluoride     189156
        Copper       199402
        Odor         178891
        Sulfate      197418
        Conductivity 163861
        Chlorine     57825
        Manganese    109583
        Total Dissolved Solids 1670
        Source       88262
        Water Temperature 168233
        Air Temperature 29728
        Month        95668
        Day          99603
        Time of Day  114519
        Target       0
        dtype: int64
```

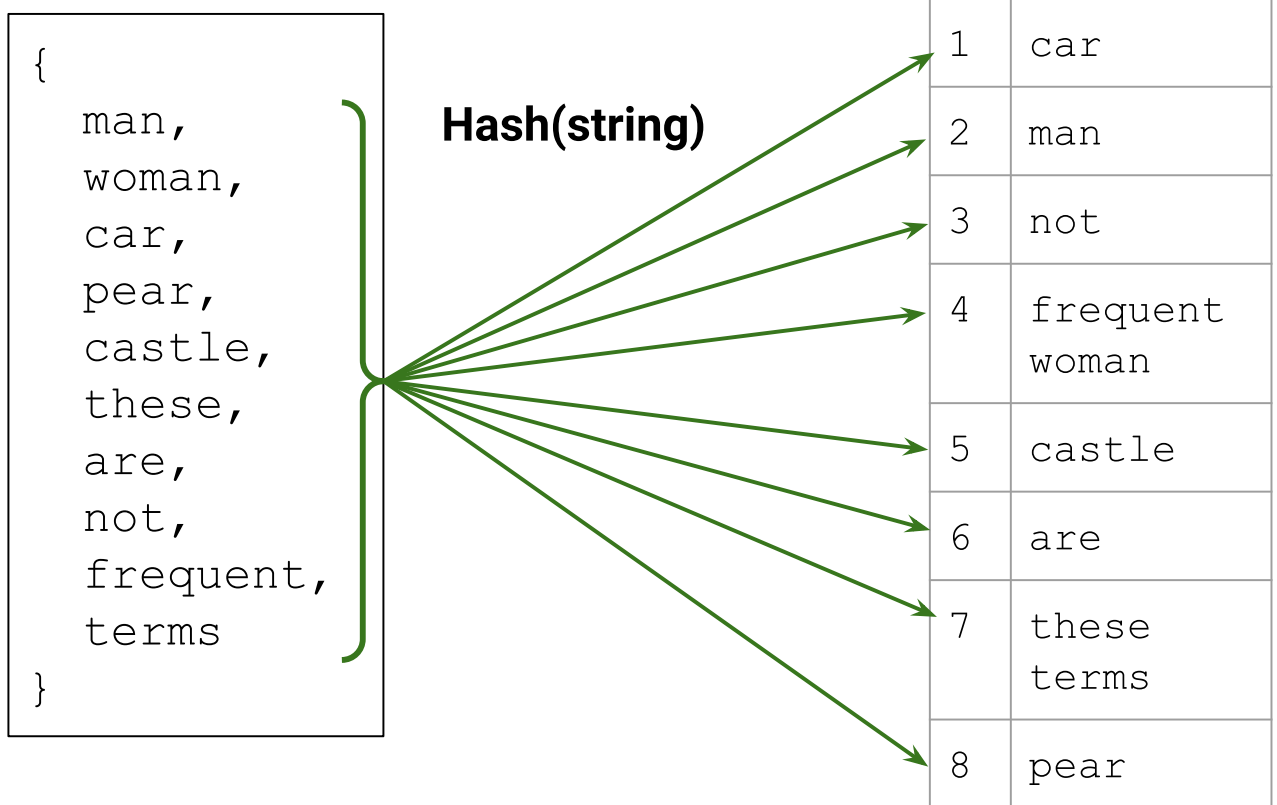
Filling all the Null Values Using Forward Fill Method

- Since there is such high standard deviation in the dataset i.e. the values are spread over wider range. So instead of filling null values by mean of a column we are using forward fill method.

```
In [10]: 1 data.fillna(method='ffill', inplace=True)
```

Converting all the Columns of String Data Type to Integer Type Using Factorization Method

- We are converting all the categorical columns into numerical data type so we can use it for training our model. Since machine learning model only accepts numerical values.



```
In [11]: 1 df_numeric = data.select_dtypes(exclude=['object'])
2 df_obj = data.select_dtypes(include=['object']).copy()
3
4 for c in df_obj:
5     df_obj[c] = pd.factorize(df_obj[c])[0]
6
7 data = pd.concat([df_obj, df_numeric], axis=1)
```

Correlation Matrix

$$r = \frac{\sum (x_i - \bar{x}) (y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$

Where,

r = Pearson Correlation Coefficient

x_i = x variable samples y_i = y variable sample

\bar{x} = mean of values in x variable \bar{y} = mean of values in y variable

In [12]:

```
1 data.corr(method='pearson').style.format("{:.3}").background_gradient(cmap=plt.get_
```

Out[12]:


	Color	Source	Month	pH	Iron	Nitrate	Chloride	Lead	Zinc
Color	1.0	0.000479	-0.000313	-0.00791	0.0425	0.0417	0.0514	0.00984	0.0205
Source	0.000479	1.0	-1.23e-05	0.000656	0.000114	0.000529	0.000452	0.000471	2.02e-05
Month	-0.000313	-1.23e-05	1.0	-0.000285	-0.000142	0.000123	8.18e-05	0.00021	-0.000125
pH	-0.00791	0.000656	-0.000285	1.0	-0.00978	-0.00946	-0.0124	-0.00232	-0.00502
Iron	0.0425	0.000114	-0.000142	-0.00978	1.0	0.0516	0.0637	0.0109	0.0252
Nitrate	0.0417	0.000529	0.000123	-0.00946	0.0516	1.0	0.0637	0.0114	0.0249
Chloride	0.0514	0.000452	8.18e-05	-0.0124	0.0637	0.0637	1.0	0.0139	0.0303
Lead	0.00984	0.000471	0.00021	-0.00232	0.0109	0.0114	0.0139	1.0	0.00582
Zinc	0.0205	2.02e-05	-0.000125	-0.00502	0.0252	0.0249	0.0303	0.00582	1.0
Turbidity	0.0564	-0.000231	-0.000425	-0.0136	0.0694	0.0703	0.0854	0.0165	0.0338
Fluoride	0.0422	0.000387	0.000147	-0.0101	0.0519	0.052	0.0636	0.0131	0.0243
Copper	0.0535	0.000304	0.000215	-0.0133	0.0673	0.0666	0.0808	0.0153	0.0323
Odor	0.0406	-0.000173	0.000335	-0.0094	0.0497	0.0494	0.0613	0.0113	0.0245
Sulfate	0.0316	-0.000307	0.000173	-0.00713	0.0404	0.0392	0.0477	0.00912	0.0192
Conductivity	-4.98e-05	6.98e-05	0.000414	0.000353	0.000269	0.000164	-0.000539	-0.000103	-3.83e-05
Chlorine	0.0369	0.000541	0.000908	-0.00774	0.045	0.0452	0.055	0.0105	0.0217
Manganese	0.0465	-0.000126	-0.000388	-0.0111	0.0588	0.0581	0.0701	0.0135	0.0272
Total Dissolved Solids	0.0233	0.000459	-0.000138	-0.00522	0.0284	0.029	0.036	0.00585	0.0138
Water Temperature	-0.000909	-0.000106	0.000282	-0.000399	0.000868	-0.00013	-0.000325	-0.000107	-0.000402
Air Temperature	-0.000538	0.000491	0.000392	0.000521	-0.000391	0.000261	-0.000247	-0.000485	2.85e-05
Day	0.000675	0.000393	-0.00596	0.000843	-9.22e-05	-6.66e-05	0.000135	0.000188	0.000218
Time of Day	-6.49e-06	0.000362	0.000222	-0.000113	0.000339	-4.43e-05	-0.000273	0.000586	0.000659
Target	0.147	-4.74e-05	-3.96e-05	-0.0347	0.181	0.183	0.223	0.042	0.0884

In [13]:

```
1 df_correlation = data.corr()
```

```
In [14]: 1 plt.rcParams['figure.figsize'] = [15, 10]
2 (df_correlation
3     .Target
4     .drop('Target')
5     .sort_values(ascending=False)
6     .plot
7     .barh())
```

Out[14]: <AxesSubplot:>




```
In [15]: 1 data.corr()['Target'].sort_values()
```

```
Out[15]: pH -0.034737
Time of Day -0.000232
Water Temperature -0.000113
Air Temperature -0.000048
Source -0.000047
Month -0.000040
Conductivity 0.000078
Day 0.000241
Lead 0.042007
Zinc 0.088429
Total Dissolved Solids 0.100231
Sulfate 0.138839
Color 0.146770
Chlorine 0.159533
Odor 0.175152
Iron 0.180897
Nitrate 0.183399
Fluoride 0.184089
Manganese 0.202702
Chloride 0.223133
Copper 0.234398
Turbidity 0.244534
Target 1.000000
Name: Target, dtype: float64
```

```
In [16]: 1 cor = data.corr()['Target'].sort_values()
```

```
In [17]: 1 df = data.copy()
```

Removing all the Columns from the Dataset whose Correlation Value is Less than 0.01 with the Target Column

- Since columns with such low correlation with not be useful for predicting the target column. So we will be dropping through column while training the predictive model

```
In [18]: 1 arr = []
2 for k, v in cor.items():
3     if abs(v) < 0.01:
4         arr.append(k)
```

```
In [19]: 1 df = df.drop(arr, axis=1)
```

```
In [20]: 1 df.head(5)
```


Out[20]:

	Color	pH	Iron	Nitrate	Chloride	Lead	Zinc	Turbidity	Fluoride	Copper
Index										
0	0	8.332988	0.000083	8.605777	122.799772	3.713298e-52	3.434827	0.022683	0.607283	0.144599
1	1	6.917863	0.000081	3.734167	227.029851	7.849262e-94	1.245317	0.019007	0.622874	0.437835
2	2	5.443762	0.020106	3.816994	230.995630	5.286616e-76	0.528280	0.319956	0.423423	0.431588
3	3	7.955339	0.143988	8.224944	178.129940	3.997118e-176	4.027879	0.166319	0.208454	0.239451
4	2	8.091909	0.002167	9.925788	186.540872	4.171069e-132	3.807511	0.004867	0.222912	0.616574

BoxPlot (Pictorial View of all the Outliers in the Data)


```
In [21]: 1 sns.set(rc={'figure.figsize':(40,60)})
          2 sns.boxplot(data=df, orient="h", palette="Set1")
```

```
Out[21]: <AxesSubplot:>
```



Removing all the Outliers from the Dataset

- Since outliers can skew the results of the predictive model. It is better to remove those from the dataset.



```
In [22]: 1 Q1 = df.quantile(0.25)
          2 Q3 = df.quantile(0.75)
          3 IQR = Q3 - Q1
```

```
In [23]: 1 ((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).sum()
```

```
Out[23]: Color          0
pH          157706
Iron         955610
Nitrate      196950
Chloride     193937
Lead        1468591
Zinc         184256
Turbidity    464338
Fluoride     207346
Copper       318721
Odor         0
Sulfate      151300
Chlorine     101584
Manganese    954536
Total Dissolved Solids  0
Target       0
dtype: int64
```

```
In [24]: 1 df = df[~((df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))).any(axis=1)]
          2 df.shape
```

```
Out[24]: (2698298, 16)
```

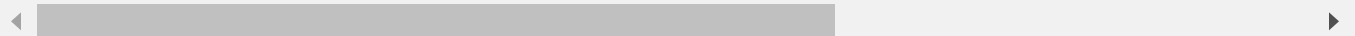
```
In [25]: 1 data.shape
```

```
Out[25]: (5956842, 23)
```

```
In [26]: 1 df.head(5)
```

```
Out[26]:
```



	Color	pH	Iron	Nitrate	Chloride	Lead	Zinc	Turbidity	Fluoride	Copper
Index										
0	0	8.332988	8.347252e-05	8.605777	122.799772	3.713298e-52	3.434827	0.022683	0.607283	0.144599
1	1	6.917863	8.053827e-05	3.734167	227.029851	7.849262e-94	1.245317	0.019007	0.622874	0.437835
4	2	8.091909	2.167128e-03	9.925788	186.540872	4.171069e-132	3.807511	0.004867	0.222912	0.616574
6	2	8.132455	5.526229e-02	4.288010	94.993978	2.919909e-52	1.770221	0.021703	1.111893	0.247116
7	0	7.258203	6.107130e-09	9.261676	182.242341	4.399852e-224	0.416478	0.047803	1.016196	0.298093



```
In [27]: 1 x = df.iloc[:, :-1].values  
2 y = df.iloc[:, -1].values
```



Standardizing all The Columns which will be used for Training the Model

- Standardizing the columns so some of the machine learning model which assign weights to each column while training should not provide higher weight to a column just based on the magnitude of their value.



Standardisation

Standardisation



comparable distributions
(m = 0.0, s = 1.0)

```
In [28]: 1 sc_X = StandardScaler()
         2 x = sc_X.fit_transform(x)
```

```
In [29]: 1 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, random_s
```

Linear Regression

```
In [30]: 1 regressor = LogisticRegression()
         2 regressor.fit(x_train, y_train)
```

```
Out[30]: LogisticRegression()
```

```
In [31]: 1 pred = regressor.predict(x_test)
```

```
In [32]: 1 accuracy = regressor.score(x_test, y_test)
2 print(f'Accuracy of Logistic Regression Model is {round(accuracy * 100, 2)} %')
```

Accuracy of Logistic Regression Model is 90.47 %

XG Boost

```
In [33]: 1 model = XGBClassifier()
```

```
In [34]: 1 model.fit(x_train, y_train)
```

```
Out[34]: XGBClassifier(base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=None, early_stopping_rounds=None,
                        enable_categorical=False, eval_metric=None, feature_types=None,
                        gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
                        interaction_constraints=None, learning_rate=None, max_bin=None,
                        max_cat_threshold=None, max_cat_to_onehot=None,
                        max_delta_step=None, max_depth=None, max_leaves=None,
                        min_child_weight=None, missing=nan, monotone_constraints=None,
                        n_estimators=100, n_jobs=None, num_parallel_tree=None,
                        predictor=None, random_state=None, ...)
```

```
In [35]: 1 pred = model.predict(x_test)
```

```
In [36]: 1 accuracy_xgb = model.score(x_test, y_test)
2 print(f'Accuracy of XG Boost Model is {round(accuracy_xgb * 100, 4)} %')
```

Accuracy of XG Boost Model is 94.7163 %

Light GBM

```
In [37]: 1 model_lgbm = LGBMClassifier()
```

```
In [38]: 1 model_lgbm.fit(x_train, y_train)
```

```
Out[38]: LGBMClassifier()
```

```
In [39]: 1 pred_new = model_lgbm.predict(x_test)
```



```
In [40]: 1 accuracy_lgbm = model_lgbm.score(x_test, y_test)
        2 print(f'Accuracy of Light Gradient Boost Model is {round(accuracy_lgbm * 100, 4)} %')
```

Accuracy of Light Gradient Boost Model is 95.003 %

CAT Boost

```
In [41]: 1 model_cat = CatBoostClassifier(verbose=0, n_estimators=100)
```

```
In [42]: 1 model_cat.fit(x_train, y_train)
```

```
Out[42]: <catboost.core.CatBoostClassifier at 0x1c113dfe100>
```

```
In [43]: 1 predictions = model_cat.predict(x_test)
```


```
In [44]: 1 accuracy_cbm = model_cat.score(x_test, y_test)
        2 print(f'Accuracy of Light Gradient Boost Model is {round(accuracy_cbm * 100, 4)} %')
```

Accuracy of Light Gradient Boost Model is 94.7898 %

Light GBM is giving the best accuracy amongst all the models

In [51]:

```
1 x = ['Linear Regression', 'XG Boost', 'Light GBM', 'CAT Boost']
2 y = [accuracy, accuracy_xgb, accuracy_lgbm, accuracy_cbm]
3 colors = ['red', 'green', 'blue', 'orange']
4 plt.rcParams['figure.figsize'] = [8,4]
5 plt.bar(x, y, color=colors)
6 plt.show()
```



Run Below Four Cells after Running all the Imports to get the Accuracy from the Final Selected Model

```
In [46]: 1 def preprocessing(dataset_path):
2
3     data = pd.read_csv(dataset_path , index_col = [0])
4     data.fillna(method='ffill', inplace=True)
5
6     df_numeric = data.select_dtypes(exclude=['object'])
7     df_obj = data.select_dtypes(include=['object']).copy()
8
9     for c in df_obj:
10         df_obj[c] = pd.factorize(df_obj[c])[0]
11
12     data = pd.concat([df_obj,df_numeric], axis=1)
13
14     cor = data.corr()['Target'].sort_values()
15     df = data.copy()
16     arr = []
17     for k, v in cor.items():
18         if abs(v) < 0.01:
19             arr.append(k)
20     df = df.drop(arr, axis=1)
21
22     Q1 = df.quantile(0.25)
23     Q3 = df.quantile(0.75)
24     IQR = Q3 - Q1
25     df = df[~((df < (Q1 - 1.5 * IQR)) |(df > (Q3 + 1.5 * IQR))).any(axis=1)]
26
27     x = df.iloc[:, :-1].values
28     y = df.iloc[:, -1].values
29     sc_X = StandardScaler()
30     x = sc_X.fit_transform(x)
31
32     return x, y
```

```
In [47]: 1 x, y = preprocessing('dataset.csv')
```

```
In [48]: 1 def training(x, y):
2
3     x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3, rand
4     model = LGBMClassifier()
5     model.fit(x_train, y_train)
6     pred = model.predict(x_test)
7     accuracy = model.score(x_test, y_test)
8     f1 = f1_score(y_test, pred)
9     return f1, accuracy
```

```
In [49]: 1 F1, Accuracy = training(x, y)
```

We are able to achieve 95.003 % Accuracy in this problem statement

```
In [50]: 1 print(f'Accuracy of the Model is {round(Accuracy * 100, 4)} %')
```

Accuracy of the Model is 95.003 %