



IntelTensorflow API Installation

```
In [1]: 1 pip install tensorflow-intel
```

```
Requirement already satisfied: h5py>=2.9.0 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (3.7.0)
Requirement already satisfied: astunparse>=1.6.0 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (1.6.3)
Requirement already satisfied: libclang>=13.0.0 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (15.0.6.1)
Requirement already satisfied: gast<=0.4.0,>=0.2.1 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (0.4.0)
Requirement already satisfied: packaging in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (21.3)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (0.29.0)
Requirement already satisfied: numpy>=1.20 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (1.21.0)
Collecting tensorboard<2.12,>=2.11
  Using cached tensorboard-2.11.2-py3-none-any.whl (6.0 MB)
Requirement already satisfied: typing-extensions>=3.6.6 in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (from tensorflow-intel) (4.3.0)
Collecting tensorflow-estimator<2.12,>=2.11.0
  Using cached tensorflow_estimator-2.11.0-py2.py3-none-any.whl (439 kB)
```

```
In [1]: 1 pip install --upgrade intel-extension-for-tensorflow[cpu]
```

```
Requirement already satisfied: intel-extension-for-tensorflow[cpu] in c:\users\ashutoshvmaadmin\anaconda3\lib\site-packages (0.0.0.dev1)
Note: you may need to restart the kernel to use updated packages.
```

```
WARNING: intel-extension-for-tensorflow 0.0.0.dev1 does not provide the extra 'cpu'
```

```
In [2]: 1 import tensorflow as tf
2 print(tf.__version__)
```

```
2.11.0
```

intel/intel-extension-for-tensorflow



Intel® Extension for TensorFlow*

Importing all the libraries

In [3]:

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 import os
5 import shutil
6 import glob
7 import cv2
8 import numpy as np
9 import math
10 import datetime
11 import time
12 import random
13 import gc
14 from tqdm import tqdm
```

In [4]:

```
1 import keras
2 import tensorflow as tf
3 from keras.models import Sequential
4 from tensorflow.keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, Max
5 from keras.preprocessing.image import ImageDataGenerator
6 from keras.optimizers import Adam
7 from sklearn.metrics import classification_report, confusion_matrix
8 from tensorflow.keras import optimizers
9 from tensorflow.keras import applications
10 from sklearn.metrics import confusion_matrix, classification_report
11 from sklearn.datasets import make_circles
12 from sklearn.metrics import accuracy_score
13 from sklearn.metrics import precision_score
14 from sklearn.metrics import recall_score
15 from sklearn.metrics import f1_score
16 from sklearn.metrics import cohen_kappa_score
17 from sklearn.metrics import roc_auc_score
18 from sklearn.model_selection import train_test_split
```

```
In [5]: 1 print(tf.__version__)
2 print(tf.config.list_physical_devices())
```

```
2.11.0
[PhysicalDevice(name='/physical_device:CPU:0', device_type='CPU')]
```

- Dividing the dataset based on labels. `train_0` contains all the images which are not infected with weed and `train_1` contains all the images which are infected with weed and likewise for test data also.

```
In [6]: 1 TRAIN_DIR = 'data/train' #Training Dataset Path
2 TEST_DIR = 'data/test'   #Testing Dataset Path
3
4 train_0 = [TRAIN_DIR + '/0' + '/' + '{}'.format(i)
5           for i in os.listdir(TRAIN_DIR + '/0')]
6 train_1 = [TRAIN_DIR + '/1' + '/' + '{}'.format(i)
7           for i in os.listdir(TRAIN_DIR + '/1')]
8
9
10 train_imgs = train_0 + train_1
11 #random.shuffle(train_imgs) # shuffle it randomly
```

```
In [7]: 1 #Getting Test Data -
2 test_0 = [TEST_DIR + '/0' + '/' + '{}'.format(i)
3           for i in os.listdir(TEST_DIR + '/0')]
4 test_1 = [TEST_DIR + '/1' + '/' + '{}'.format(i)
5           for i in os.listdir(TEST_DIR + '/1')]
6
7
8 test_imgs = test_0 + test_1
9 random.shuffle(test_imgs)
```

- function to read and process the images to an acceptable format for our model

```
In [8]: 1 nrows = 224
2 ncolumns = 224
3
4 #A function to read and process the images to an acceptable format for our model
5 def read_and_process_image(list_of_images):
6     """
7         Returns two arrays:
8             X is an array of resized images
9             y is an array of labels
10        """
11
12     X = [] # images
13     y = [] # labels
14
15     for image in tqdm(list_of_images):
16         X.append(cv2.resize(cv2.imread(image, cv2.IMREAD_COLOR),
17                             (nrows,ncolumns), interpolation=cv2.INTER_CUBIC)) #Read
18         #get the labels
19         if '0' in image[:12]:
20             y.append(0)
21         elif '1' in image[:12]:
22             y.append(1)
23
24     return X, y
```

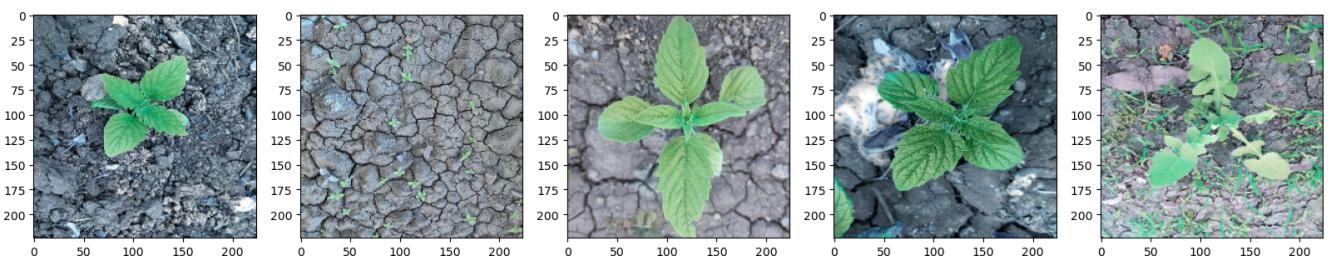
Reading images and labels from the training dataset

```
In [9]: 1 X, y = read_and_process_image(train_imgs)
```

100%|██████████| 10
39/1039 [00:07<00:00, 131.79it/s]

```
In [10]: 1 plt.figure(figsize=(20,10))
2 columns = 5
3 for i in tqdm(range(columns)):
4     plt.subplot(5 // columns + 1, columns, i + 1)
5     plt.imshow(X[i])
```

100%|██████████| 5/5 [00:00<00:00, 46.76it/s]

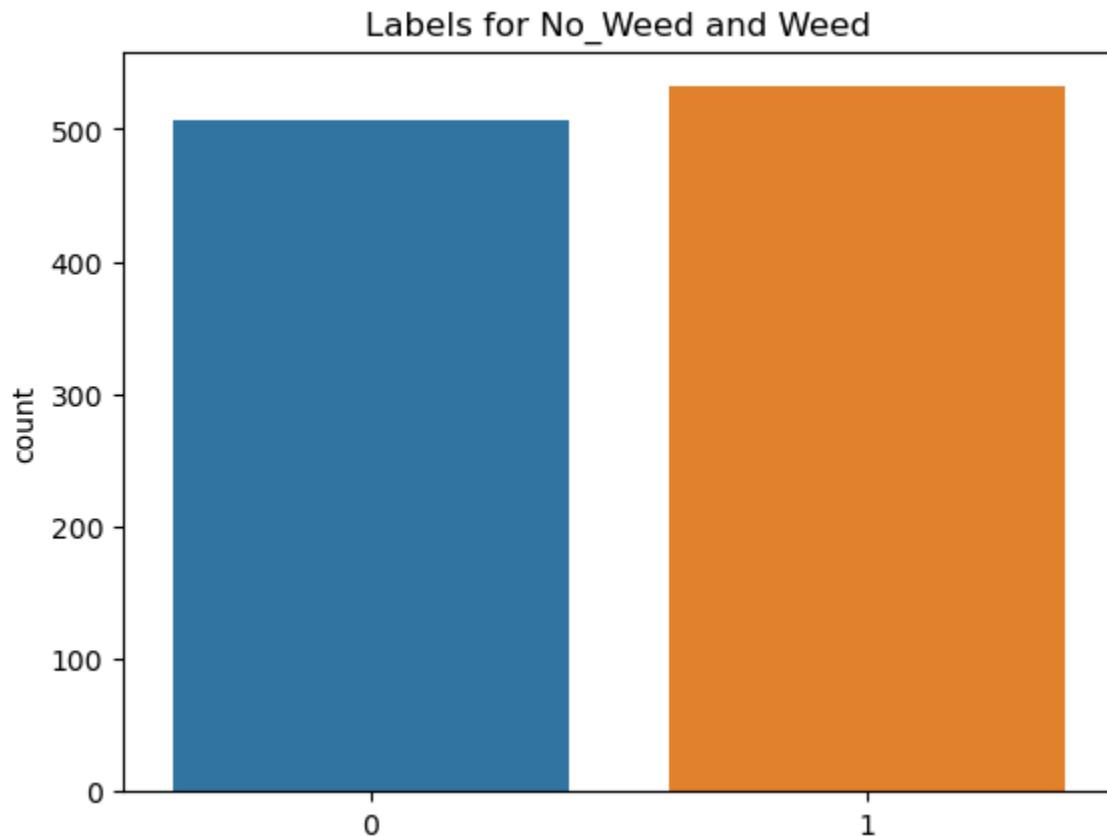


In [11]:

```
1 del train_imgs
2 gc.collect()
3
4 #Convert list to numpy array
5 X = np.array(X)
6 y = np.array(y)
7
8
9 #Lets plot the label to be sure we just have two class
10 sns.countplot(y)
11 plt.title('Labels for No_Weed and Weed')
```

C:\Users\ashutoshvmadmin\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

Out[11]: Text(0.5, 1.0, 'Labels for No_Weed and Weed')

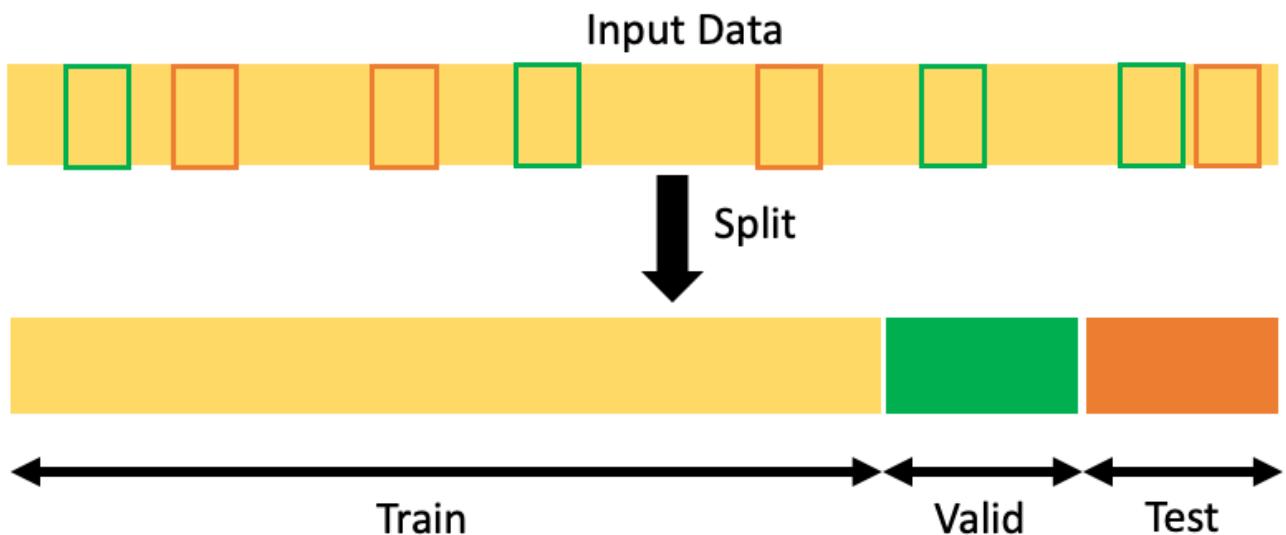


```
1 <hr style = "border-top: 3px solid black" >
```

Creating Training and Testing Data

```
1 X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.20,
random_state=2)
2
3 print("Shape of train images is:", X_train.shape)
4 print("Shape of validation images is:", X_val.shape)
5 print("Shape of labels is:", y_train.shape)
```

```
6 print("Shape of labels is:", y_val.shape)
```



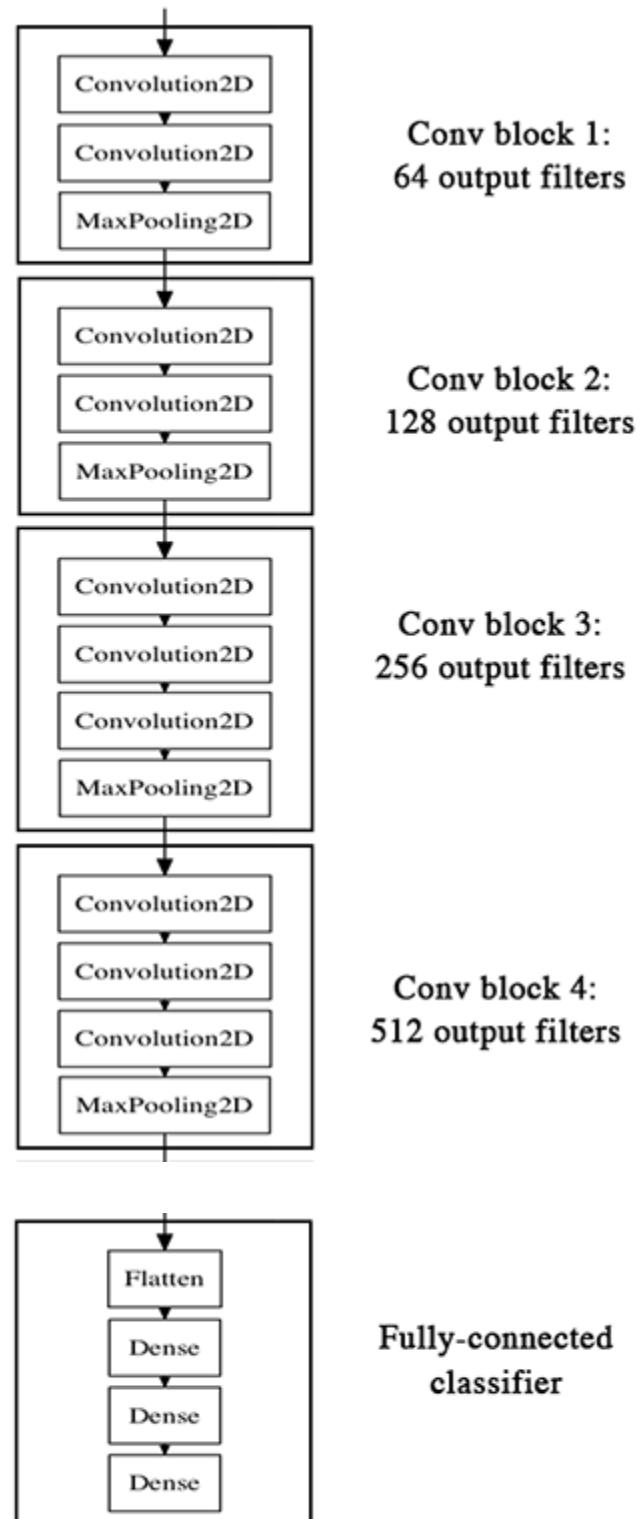
```
In [13]: 1 print("Shape of train images is:", X.shape)
          2 print("Shape of labels is:", y.shape)
```

Shape of train images is: (1039, 224, 224, 3)
Shape of labels is: (1039,)

```
In [14]: 1 del X
          2 del y
          3 gc.collect()
          4
          5 #get the length of the train and validation data
          6 ntrain = len(X_train)
          7 nval = len(X_val)
          8 print(ntrain)
          9 print(nval)
          10
          11 #We will use a batch size of 32. Note: batch size should be a factor of 2.***4,8,16.
          12 batch_size = 32
```

831
208

Preparing our Initial Model for Training



```
In [35]: 1 model = Sequential()
2
3 model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224, 224, 3)))
4 model.add(MaxPooling2D((2, 2)))
5
6 model.add(Conv2D(64, (3, 3), activation='relu'))
7 model.add(MaxPooling2D((2, 2)))
8
9 model.add(Conv2D(128, (3, 3), activation='relu'))
10 model.add(MaxPooling2D((2, 2)))
11
12 model.add(Conv2D(256, (3, 3), activation='relu'))
13 model.add(MaxPooling2D((2, 2)))
14
15 model.add(Flatten())
16 model.add(Dropout(0.5)) #Dropout for regularization
17 model.add(Dense(512, activation='relu'))
18 model.add(Dense(1, activation='sigmoid'))
```

```
In [16]: 1 model.compile(loss='binary_crossentropy', optimizer=optimizers.RMSprop(lr=1e-4), me
```

WARNING:absl:`lr` is deprecated, please use `learning_rate` instead, or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy.RMSprop.

```
In [17]: 1 train_datagen = ImageDataGenerator(rescale=1./255, #Scale the image between 0 and
2                                     rotation_range=50,
3                                     width_shift_range=0.2,
4                                     height_shift_range=0.2,
5                                     shear_range=0.2,
6                                     zoom_range=0.2,
7                                     horizontal_flip=True,)
8
9 #added
10 train_datagen = ImageDataGenerator(rescale=1./255 #Scale the image between 0 and .)
11
12
13
14 val_datagen = ImageDataGenerator(rescale=1./255) #We do not augment validation dat
```

```
In [18]: 1 train_generator = train_datagen.flow(X_train, y_train, batch_size=batch_size)
2 val_generator = val_datagen.flow(X_val, y_val, batch_size=batch_size)
```

Training our Initial Model

```
In [19]: 1 history = model.fit_generator(train_generator,
2                                steps_per_epoch=ntrain // batch_size,
3                                epochs=50,
4                                validation_data=val_generator,
5                                validation_steps=nval // batch_size)
```

C:\Users\ashutoshv\appData\Local\Temp\2\ipykernel_1980\1008497584.py:1: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

```
history = model.fit_generator(train_generator,
```

Epoch 1/50
25/25 [=====] - 24s 894ms/step - loss: 1.1702 - acc: 0.4981 - val_loss: 0.6895 - val_acc: 0.5208
Epoch 2/50
25/25 [=====] - 22s 895ms/step - loss: 0.6942 - acc: 0.5019 - val_loss: 0.6872 - val_acc: 0.5365
Epoch 3/50
25/25 [=====] - 23s 905ms/step - loss: 0.7580 - acc: 0.5257 - val_loss: 0.6935 - val_acc: 0.4792
Epoch 4/50
25/25 [=====] - 23s 938ms/step - loss: 0.6936 - acc: 0.5820 - val_loss: 0.7331 - val_acc: 0.5312
Epoch 5/50
25/25 [=====] - 23s 938ms/step - loss: 0.6754 - acc: 0.5970 - val_loss: 0.5516 - val_acc: 0.8750
Epoch 6/50
25/25 [=====] - 23s 915ms/step - loss: 0.6581 - acc: 0.7259 - val_loss: 0.4477 - val_acc: 0.8490
Epoch 7/50
25/25 [=====] - 22s 896ms/step - loss: 0.4954 - acc: 0.7950 - val_loss: 0.5159 - val_acc: 0.7865
Epoch 8/50
25/25 [=====] - 23s 910ms/step - loss: 0.4413 - acc: 0.8198 - val_loss: 0.9165 - val_acc: 0.5469
Epoch 9/50
25/25 [=====] - 23s 922ms/step - loss: 0.4941 - acc: 0.7822 - val_loss: 1.0834 - val_acc: 0.6510
Epoch 10/50
25/25 [=====] - 22s 901ms/step - loss: 0.4391 - acc: 0.8436 - val_loss: 0.4376 - val_acc: 0.8333
Epoch 11/50
25/25 [=====] - 22s 889ms/step - loss: 0.4226 - acc: 0.8123 - val_loss: 0.4079 - val_acc: 0.9167
Epoch 12/50
25/25 [=====] - 23s 912ms/step - loss: 0.3645 - acc: 0.8548 - val_loss: 0.3276 - val_acc: 0.9115
Epoch 13/50
25/25 [=====] - 22s 893ms/step - loss: 0.3661 - acc: 0.8573 - val_loss: 0.4272 - val_acc: 0.8438
Epoch 14/50
25/25 [=====] - 22s 900ms/step - loss: 0.3355 - acc: 0.8561 - val_loss: 0.5437 - val_acc: 0.7031
Epoch 15/50
25/25 [=====] - 22s 897ms/step - loss: 0.3314 - acc: 0.8611 - val_loss: 0.3131 - val_acc: 0.9010
Epoch 16/50
25/25 [=====] - 22s 898ms/step - loss: 0.3315 - acc: 0.8636 - val_loss: 0.2692 - val_acc: 0.9219
Epoch 17/50
25/25 [=====] - 23s 901ms/step - loss: 0.2728 - acc: 0.8886 - val_loss: 0.4138 - val_acc: 0.8490
Epoch 18/50
25/25 [=====] - 22s 897ms/step - loss: 0.3651 - acc: 0.8498 - val_loss: 0.2828 - val_acc: 0.9271
Epoch 19/50
25/25 [=====] - 23s 903ms/step - loss: 0.2738 - acc: 0.8874 - val_loss: 0.7057 - val_acc: 0.7812
Epoch 20/50
25/25 [=====] - 23s 907ms/step - loss: 0.2545 - acc: 0.8861 -

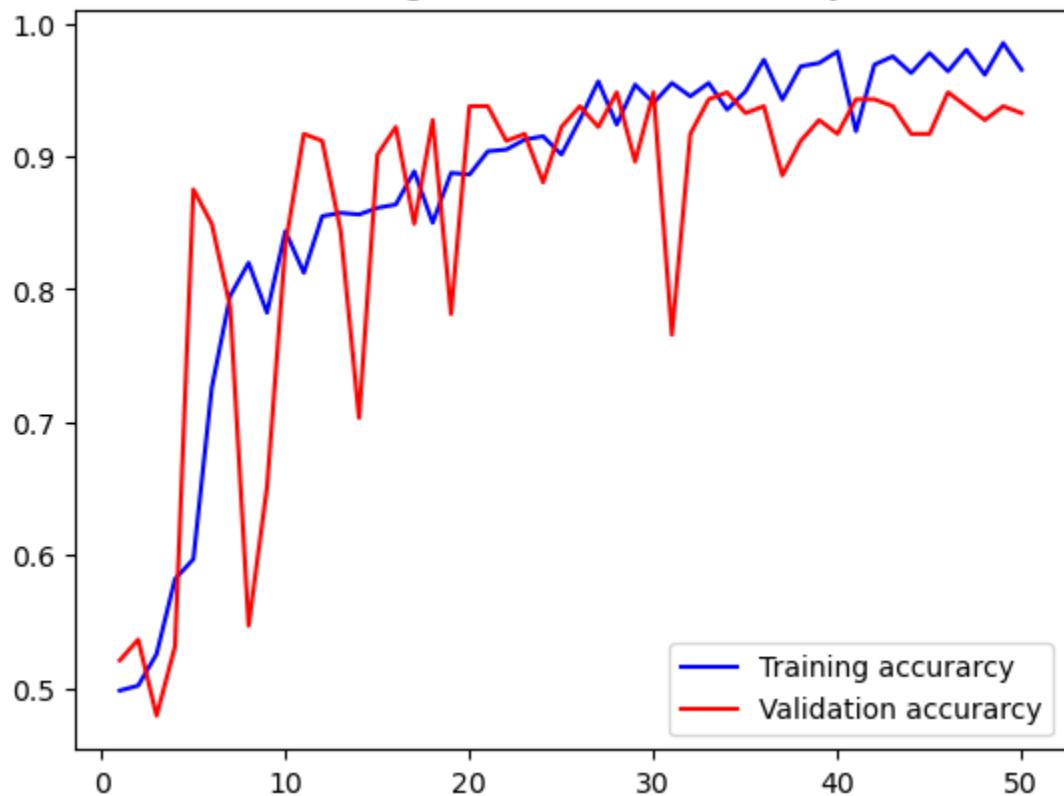
```
val_loss: 0.2847 - val_acc: 0.9375
Epoch 21/50
25/25 [=====] - 22s 888ms/step - loss: 0.2434 - acc: 0.9036 -
val_loss: 0.2319 - val_acc: 0.9375
Epoch 22/50
25/25 [=====] - 22s 894ms/step - loss: 0.2473 - acc: 0.9049 -
val_loss: 0.2878 - val_acc: 0.9115
Epoch 23/50
25/25 [=====] - 22s 890ms/step - loss: 0.2464 - acc: 0.9124 -
val_loss: 0.2719 - val_acc: 0.9167
Epoch 24/50
25/25 [=====] - 22s 898ms/step - loss: 0.2397 - acc: 0.9149 -
val_loss: 0.3673 - val_acc: 0.8802
Epoch 25/50
25/25 [=====] - 22s 896ms/step - loss: 0.2672 - acc: 0.9011 -
val_loss: 0.2875 - val_acc: 0.9219
Epoch 26/50
25/25 [=====] - 22s 895ms/step - loss: 0.2159 - acc: 0.9274 -
val_loss: 0.2535 - val_acc: 0.9375
Epoch 27/50
25/25 [=====] - 22s 880ms/step - loss: 0.1515 - acc: 0.9562 -
val_loss: 0.2762 - val_acc: 0.9219
Epoch 28/50
25/25 [=====] - 22s 888ms/step - loss: 0.2218 - acc: 0.9237 -
val_loss: 0.3218 - val_acc: 0.9479
Epoch 29/50
25/25 [=====] - 22s 889ms/step - loss: 0.1576 - acc: 0.9537 -
val_loss: 0.3354 - val_acc: 0.8958
Epoch 30/50
25/25 [=====] - 22s 890ms/step - loss: 0.1887 - acc: 0.9399 -
val_loss: 0.3288 - val_acc: 0.9479
Epoch 31/50
25/25 [=====] - 22s 886ms/step - loss: 0.1362 - acc: 0.9549 -
val_loss: 1.4850 - val_acc: 0.7656
Epoch 32/50
25/25 [=====] - 22s 872ms/step - loss: 0.1870 - acc: 0.9449 -
val_loss: 0.3157 - val_acc: 0.9167
Epoch 33/50
25/25 [=====] - 22s 885ms/step - loss: 0.1226 - acc: 0.9549 -
val_loss: 0.3075 - val_acc: 0.9427
Epoch 34/50
25/25 [=====] - 22s 883ms/step - loss: 0.1659 - acc: 0.9349 -
val_loss: 0.2430 - val_acc: 0.9479
Epoch 35/50
25/25 [=====] - 22s 895ms/step - loss: 0.1308 - acc: 0.9487 -
val_loss: 0.3297 - val_acc: 0.9323
Epoch 36/50
25/25 [=====] - 22s 888ms/step - loss: 0.0764 - acc: 0.9725 -
val_loss: 0.3521 - val_acc: 0.9375
Epoch 37/50
25/25 [=====] - 22s 884ms/step - loss: 0.1460 - acc: 0.9424 -
val_loss: 0.4456 - val_acc: 0.8854
Epoch 38/50
25/25 [=====] - 23s 912ms/step - loss: 0.0902 - acc: 0.9675 -
val_loss: 0.4441 - val_acc: 0.9115
Epoch 39/50
25/25 [=====] - 22s 899ms/step - loss: 0.0843 - acc: 0.9700 -
val_loss: 0.3827 - val_acc: 0.9271
Epoch 40/50
```

```
25/25 [=====] - 22s 889ms/step - loss: 0.0724 - acc: 0.9787 -  
val_loss: 0.4470 - val_acc: 0.9167  
Epoch 41/50  
25/25 [=====] - 22s 894ms/step - loss: 0.1918 - acc: 0.9186 -  
val_loss: 0.2569 - val_acc: 0.9427  
Epoch 42/50  
25/25 [=====] - 22s 880ms/step - loss: 0.0720 - acc: 0.9687 -  
val_loss: 0.3408 - val_acc: 0.9427  
Epoch 43/50  
25/25 [=====] - 22s 892ms/step - loss: 0.0737 - acc: 0.9750 -  
val_loss: 0.3546 - val_acc: 0.9375  
Epoch 44/50  
25/25 [=====] - 22s 896ms/step - loss: 0.0974 - acc: 0.9625 -  
val_loss: 0.4139 - val_acc: 0.9167  
Epoch 45/50  
25/25 [=====] - 22s 892ms/step - loss: 0.0483 - acc: 0.9775 -  
val_loss: 0.4449 - val_acc: 0.9167  
Epoch 46/50  
25/25 [=====] - 23s 939ms/step - loss: 0.0934 - acc: 0.9637 -  
val_loss: 0.3536 - val_acc: 0.9479  
Epoch 47/50  
25/25 [=====] - 23s 904ms/step - loss: 0.0517 - acc: 0.9800 -  
val_loss: 0.3420 - val_acc: 0.9375  
Epoch 48/50  
25/25 [=====] - 22s 892ms/step - loss: 0.1289 - acc: 0.9612 -  
val_loss: 0.4328 - val_acc: 0.9271  
Epoch 49/50  
25/25 [=====] - 22s 890ms/step - loss: 0.0334 - acc: 0.9850 -  
val_loss: 0.4142 - val_acc: 0.9375  
Epoch 50/50  
25/25 [=====] - 22s 890ms/step - loss: 0.0738 - acc: 0.9650 -  
val_loss: 0.2982 - val_acc: 0.9323
```

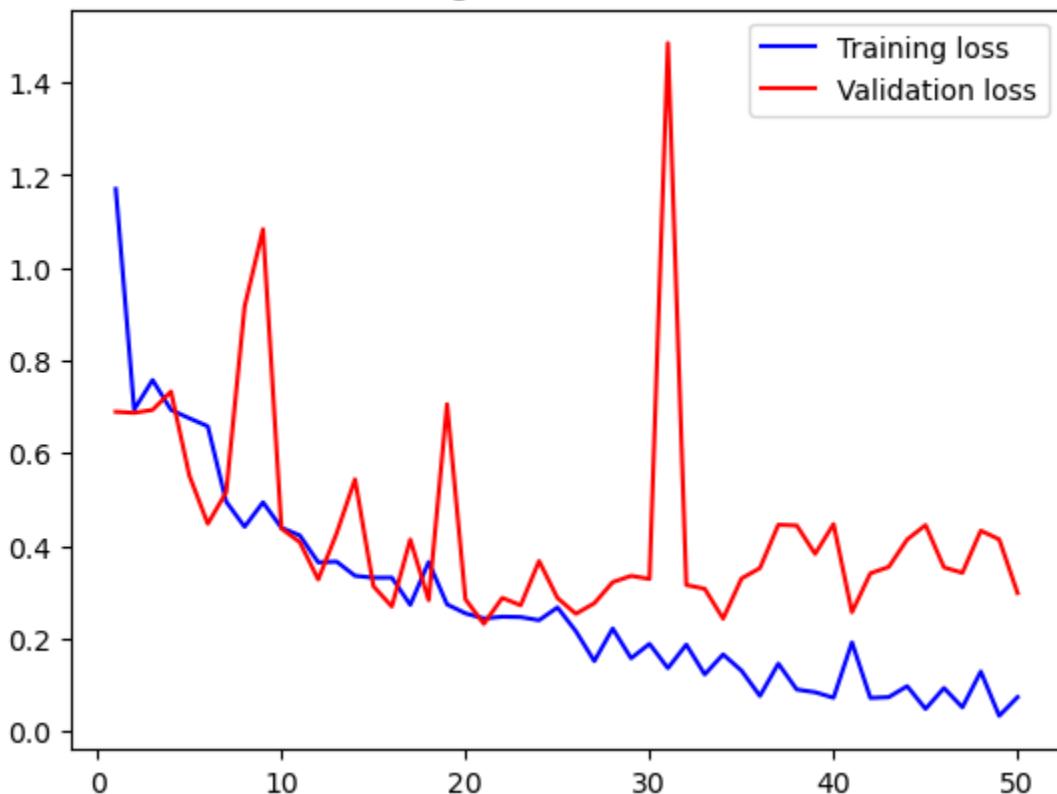
In [21]:

```
1 acc = history.history['acc']
2 val_acc = history.history['val_acc']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5
6 epochs = range(1, len(acc) + 1)
7
8 #Train and validation accuracy
9 plt.plot(epochs, acc, 'b', label='Training accuracy')
10 plt.plot(epochs, val_acc, 'r', label='Validation accuracy')
11 plt.title('Training and Validation accuracy')
12 plt.legend()
13
14
15
16 plt.figure()
17 #Train and validation loss
18 plt.plot(epochs, loss, 'b', label='Training loss')
19 plt.plot(epochs, val_loss, 'r', label='Validation loss')
20 plt.title('Training and Validation loss')
21 plt.legend()
22
23 plt.show()
```

Training and Validation accuracy



Training and Validation loss



```
In [22]: 1 X_test, y_test = read_and_process_image(test_imgs) #Y_test in this case will be empty
2 x = np.array(X_test)
3 test_datagen = ImageDataGenerator(rescale=1./255)

100%|██████████| 261/261 [00:03<00:00, 76.21it/s]
```

Validating our Initial Model

```
In [23]: 1 out = model.evaluate_generator(val_generator)
2 print(out)
```

```
C:\Users\ashutoshvmadmin\AppData\Local\Temp\2\ipykernel_1980\4193496103.py:1: UserWarning: `Model.evaluate_generator` is deprecated and will be removed in a future version.
Please use `Model.evaluate`, which supports generators.
out = model.evaluate_generator(val_generator)
[0.32266777753829956, 0.932692289352417]
```

```
In [24]: 1 Y_pred = model.predict(X_val)
2 print(Y_pred.shape)
3 #y_pred = np.argmax(Y_pred, axis=1) - used for multiclass
4 y_pred = (Y_pred > 0.5) * 1.0
5 y_pred = y_pred.reshape(y_val.shape)
6 y_pred.sum()
```

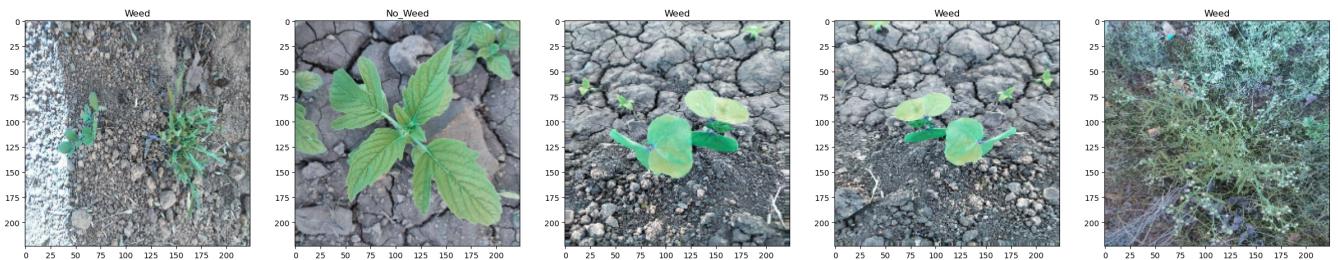
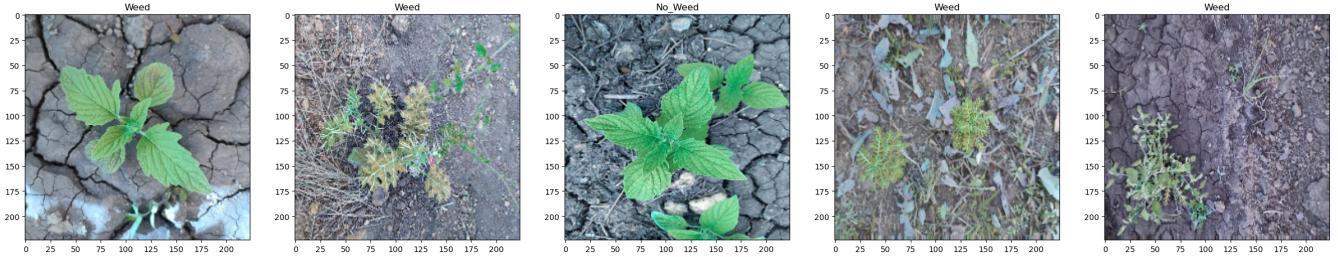
```
7/7 [=====] - 2s 194ms/step
(208, 1)
```

```
Out[24]: 104.0
```

In [25]:

```
1 i = 0
2 text_labels = []
3 plt.figure(figsize=(30,20))
4
5 for batch in test_datagen.flow(x, batch_size=1):
6     pred = model.predict(batch)
7     if pred > 0.5:
8         text_labels.append('Weed')
9     else:
10        text_labels.append('No_Weed')
11 plt.subplot(5 // columns + 1, columns, i + 1)
12 plt.title(text_labels[i])
13 #print(batch[0])
14 imgplot = plt.imshow(batch[0])
15 i += 1
16 if i % 10 == 0:
17     break
18 plt.show()
```

```
1/1 [=====] - 0s 127ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 50ms/step
```



In [26]:

```
1 print('Confusion Matrix')
2 print(confusion_matrix(y_val, y_pred))
3
4
5 print('Classification Report')
6 target_names = ['Weed', 'No Weed']
7 print(classification_report(y_val, y_pred,target_names=target_names))
8
9 # accuracy:  $(tp + tn) / (p + n)$ 
10 accuracy = accuracy_score(y_val, y_pred)
11 print('Accuracy: %f' % accuracy)
12 # precision  $tp / (tp + fp)$ 
13 precision = precision_score(y_val, y_pred)
14 print('Precision: %f' % precision)
15 # recall:  $tp / (tp + fn)$ 
16 recall = recall_score(y_val, y_pred)
17 print('Recall: %f' % recall)
18 # f1:  $2 tp / (2 tp + fp + fn)$ 
19 f1 = f1_score(y_val, y_pred)
20 print('F1 score: %f' % f1)
```

Confusion Matrix

```
[[ 96   3]
 [  8 101]]
```

Classification Report

	precision	recall	f1-score	support
Weed	0.92	0.97	0.95	99
No Weed	0.97	0.93	0.95	109
accuracy			0.95	208
macro avg	0.95	0.95	0.95	208
weighted avg	0.95	0.95	0.95	208

Accuracy: 0.947115

Precision: 0.971154

Recall: 0.926606

F1 score: 0.948357

		ACTUAL VALUES	
PREDICTED VALUES	Positive	Positive	Negative
	Positive	TP	FP
		FN	TN

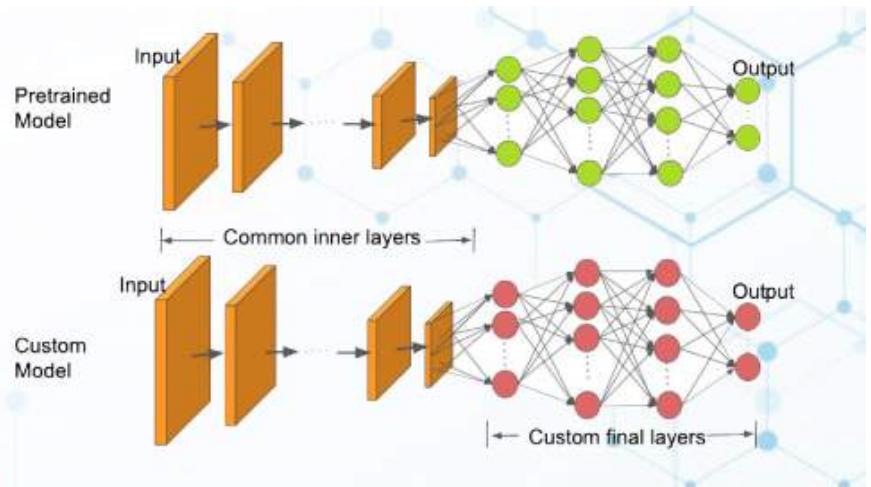
The diagram illustrates a 2x2 confusion matrix for classification performance. The columns represent 'ACTUAL VALUES' (Positive or Negative) and the rows represent 'PREDICTED VALUES' (Positive or Negative). The matrix cells are labeled: TP (True Positive, top-left, green), FP (False Positive, top-right, red), FN (False Negative, bottom-left, red), and TN (True Negative, bottom-right, green). Four annotations with arrows point to specific cells:

- An arrow points from the top-left cell (TP) to a green box: "The predicted value is positive and its positive".
- An arrow points from the top-right cell (FP) to a red box: "Type I error : The predicted value is positive but it False".
- An arrow points from the bottom-left cell (FN) to a red box: "Type II error : The predicted value is negative but its positive".
- An arrow points from the bottom-right cell (TN) to a green box: "The predicted value is Negative and its Negative".

Transfer Learning

- Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task. Transfer learning can be used when the dataset is small, by using a pre-trained model on similar images we can easily achieve high performance.

Transfer Learning



```
In [27]: 1 base_model = tf.keras.applications.MobileNetV2(input_shape = (224, 224, 3), include
```

```
In [28]: 1 base_model.trainable = False
```

```
In [29]: 1 model = tf.keras.Sequential([base_model,
 2                         tf.keras.layers.GlobalAveragePooling2D(),
 3                         tf.keras.layers.Dropout(0.2),
 4                         tf.keras.layers.Dense(1, activation="sigmoid")
 5])
```

Training our Final Model

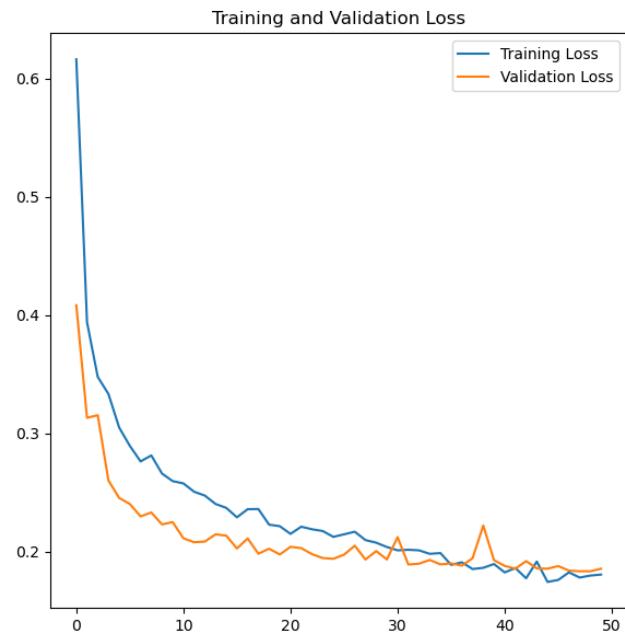
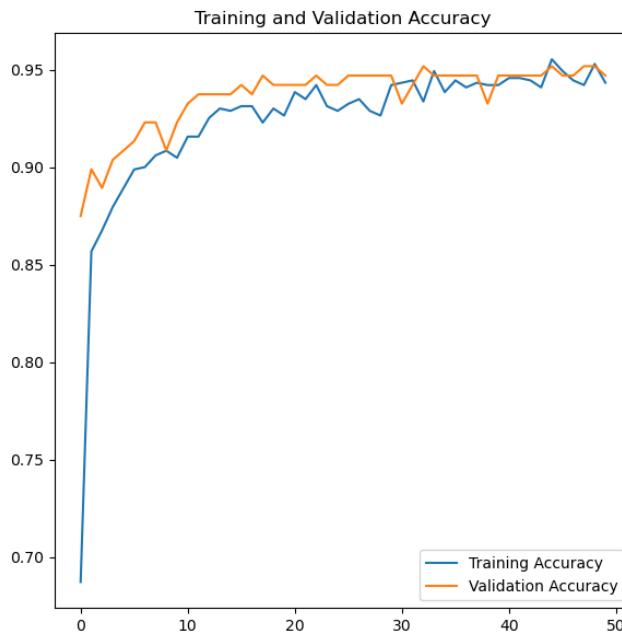
In [30]:

```
1 base_learning_rate = 0.00001
2 model.compile(optimizer=tf.keras.optimizers.Adam(lr=base_learning_rate),
3                 loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
4                 metrics=['accuracy'])
5
6 history = model.fit(X_train,y_train,epochs = 50 , validation_data = (X_val, y_val))

0.9302 - val_loss: 0.2024 - val_accuracy: 0.9423
Epoch 20/50
26/26 [=====] - 13s 489ms/step - loss: 0.2215 - accuracy:
0.9266 - val_loss: 0.1975 - val_accuracy: 0.9423
Epoch 21/50
26/26 [=====] - 13s 487ms/step - loss: 0.2150 - accuracy:
0.9386 - val_loss: 0.2040 - val_accuracy: 0.9423
Epoch 22/50
26/26 [=====] - 13s 493ms/step - loss: 0.2210 - accuracy:
0.9350 - val_loss: 0.2029 - val_accuracy: 0.9423
Epoch 23/50
26/26 [=====] - 13s 488ms/step - loss: 0.2189 - accuracy:
0.9422 - val_loss: 0.1979 - val_accuracy: 0.9471
Epoch 24/50
26/26 [=====] - 13s 493ms/step - loss: 0.2174 - accuracy:
0.9314 - val_loss: 0.1945 - val_accuracy: 0.9423
Epoch 25/50
26/26 [=====] - 13s 488ms/step - loss: 0.2124 - accuracy:
0.9290 - val_loss: 0.1939 - val_accuracy: 0.9423
Epoch 26/50
```

In [31]:

```
1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3 loss = history.history['loss']
4 val_loss = history.history['val_loss']
5 epochs_range = range(50)
6
7 plt.figure(figsize=(15, 15))
8 plt.subplot(2, 2, 1)
9 #Train and validation accuracy
10 plt.plot(epochs_range, acc, label='Training Accuracy')
11 plt.plot(epochs_range, val_acc, label='Validation Accuracy')
12 plt.legend(loc='lower right')
13 plt.title('Training and Validation Accuracy')
14
15 #Train and validation loss
16 plt.subplot(2, 2, 2)
17 plt.plot(epochs_range, loss, label='Training Loss')
18 plt.plot(epochs_range, val_loss, label='Validation Loss')
19 plt.legend(loc='upper right')
20 plt.title('Training and Validation Loss')
21 plt.show()
```



Calculating all the Evaluation Metrics

In [32]:

```
1 predictions = model.predict(X_val)
2 predictions = predictions.reshape(-1)
3
4 for i in range(len(predictions)):
5     if predictions[i] > 0.5:
6         predictions[i] = 1
7     else:
8         predictions[i] = 0
9
10 print(classification_report(y_val, predictions, target_names = ['No Weed (Class 0)']))
```

```
7/7 [=====] - 3s 363ms/step
      precision    recall   f1-score   support
No Weed (Class 0)       0.92      0.97      0.95      99
  Weed (Class 1)       0.97      0.93      0.95     109

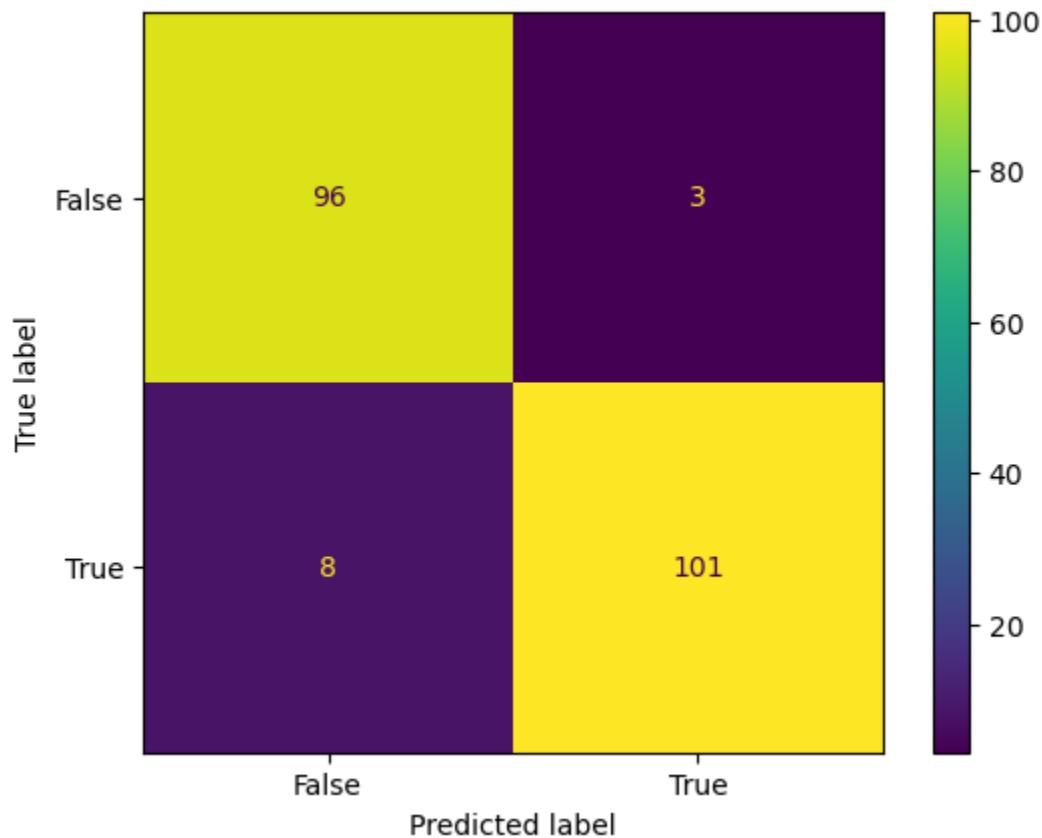
      accuracy           0.95      0.95      0.95     208
    macro avg           0.95      0.95      0.95     208
 weighted avg          0.95      0.95      0.95     208
```

Confusion Metrics

In [42]:

```
1 from sklearn import metrics
2
3 confusion_matrix = metrics.confusion_matrix(y_val, predictions)
4 cm_display = metrics.ConfusionMatrixDisplay(confusion_matrix = confusion_matrix, di
```

```
In [43]: 1 cm_display.plot()  
2 plt.show()
```



Weed Detection Assignment Summary

We are able to achieve F1 Score of 0.95 with this Model

```
In [44]: 1 f1 = f1_score(y_val, predictions)
```

```
In [45]: 1 from IPython.display import Markdown, display  
2  
3 f1_text = f"**F1 Score:** {f1:.2f}"  
4 display(Markdown(f1_text))
```

F1 Score: 0.95